

# Final Project

London Wagner, Erik Lovece, Carmen Canedo

2021-04-24

## 1 Introduction

Our final project analyzes the student performance dataset from the UCI Machine Learning Repository, originally gathered by Paulo Cortez from the University of Minho. This dataset measures the final student grade in a Portuguese class based on a variety of predictors. These predictors cover numerous aspects of not only students' academic lives, but also family life predictors such as parental employment, and personal predictors like whether or not they have home internet access and whether or not they are in a romantic relationship.

We seek to answer the question of what predictors have the greatest influence in how a student does in class. Conventional wisdom seems to dictate that high-achieving students have come from particularly favorable academic, filial, and personal environments, and previous studies have confirmed this. Our model, if properly constructed along the best machine learning practices, should corroborate this, although unexpected conclusions may also lie in store.

Our workflow for finding a sufficient model from which we will draw our conclusions is as follows:

- 1) Run a linear regression model with final grade as the response and all other variables as predictors
- 2) Use best subset, forward step, and backwards step to select variables for a reduced model
- 3) Use ridge and lasso to conduct further dimension reduction
- 4) Use cross validation methods to determine which model predicts the final grade with the greatest accuracy
- 5) Make more definitive determinations based on the chosen model.

## 2 Loading & Cleaning Data

```
student_por <- read_csv2("data/student-por.csv")
```

```
student_por
```

```
## # A tibble: 649 x 33
```

```
##   school sex   age address famsize Pstatus Medu Fedu Mjob Fjob reason
##   <chr> <chr> <dbl> <chr>   <chr>   <chr>   <dbl> <dbl> <chr>   <chr>   <chr>
## 1 GP    F     18 U      GT3     A       4     4 at_home teach~ course
## 2 GP    F     17 U      GT3     T       1     1 at_home other  course
## 3 GP    F     15 U      LE3     T       1     1 at_home other  other
## 4 GP    F     15 U      GT3     T       4     2 health servi~ home
## 5 GP    F     16 U      GT3     T       3     3 other  other  home
```

```

## 6 GP      M      16 U      LE3      T      4      3 servic~ other reputa~
## 7 GP      M      16 U      LE3      T      2      2 other  other  home
## 8 GP      F      17 U      GT3      A      4      4 other  teach~ home
## 9 GP      M      15 U      LE3      A      3      2 servic~ other  home
## 10 GP     M      15 U      GT3      T      3      4 other  other  home
## # ... with 639 more rows, and 22 more variables: guardian <chr>,
## #   traveltime <dbl>, studytime <dbl>, failures <dbl>, schoolsup <chr>,
## #   famsup <chr>, paid <chr>, activities <chr>, nursery <chr>, higher <chr>,
## #   internet <chr>, romantic <chr>, famrel <dbl>, freetime <dbl>, goout <dbl>,
## #   Dalc <dbl>, Walc <dbl>, health <dbl>, absences <dbl>, G1 <dbl>, G2 <dbl>,
## #   G3 <dbl>

```

The student attributes and grades forming the predictors and response, quoted verbatim from a text file provided with the dataset, are as follows:

- 1 school - student's school (binary: "GP" - Gabriel Pereira or "MS" - Mousinho da Silveira)
- 2 sex - student's sex (binary: "F" - female or "M" - male)
- 3 age - student's age (numeric: from 15 to 22)
- 4 address - student's home address type (binary: "U" - urban or "R" - rural)
- 5 famsize - family size (binary: "LE3" - less or equal to 3 or "GT3" - greater than 3)
- 6 Pstatus - parent's cohabitation status (binary: "T" - living together or "A" - apart)
- 7 Medu - mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)
- 8 Fedu - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)
- 9 Mjob - mother's job (nominal: "teacher", "health" care related, civil "services" (e.g. administrative or police), "at\_home" or "other")
- 10 Fjob - father's job (nominal: "teacher", "health" care related, civil "services" (e.g. administrative or police), "at\_home" or "other")
- 11 reason - reason to choose this school (nominal: close to "home", school "reputation", "course" preference or "other")
- 12 guardian - student's guardian (nominal: "mother", "father" or "other")
- 13 traveltime - home to school travel time (numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour)
- 14 studytime - weekly study time (numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours)
- 15 failures - number of past class failures (numeric: n if  $1 \leq n < 3$ , else 4)
- 16 schoolsup - extra educational support (binary: yes or no)
- 17 famsup - family educational support (binary: yes or no)
- 18 paid - extra paid classes within the course subject (Math or Portuguese) (binary: yes or no)
- 19 activities - extra-curricular activities (binary: yes or no)
- 20 nursery - attended nursery school (binary: yes or no)
- 21 higher - wants to take higher education (binary: yes or no)
- 22 internet - Internet access at home (binary: yes or no)

23 romantic - with a romantic relationship (binary: yes or no)  
 24 famrel - quality of family relationships (numeric: from 1 - very bad to 5 - excellent)  
 25 freetime - free time after school (numeric: from 1 - very low to 5 - very high)  
 26 goout - going out with friends (numeric: from 1 - very low to 5 - very high)  
 27 Dalc - workday alcohol consumption (numeric: from 1 - very low to 5 - very high)  
 28 Walc - weekend alcohol consumption (numeric: from 1 - very low to 5 - very high)  
 29 health - current health status (numeric: from 1 - very bad to 5 - very good)  
 30 absences - number of school absences (numeric: from 0 to 93)  
 31 G1 - first period grade (numeric: from 0 to 20)  
 31 G2 - second period grade (numeric: from 0 to 20)  
 32 G3 - final grade (numeric: from 0 to 20, output target)

```
student_por <-
  student_por %>%
  mutate(school = factor(school),
         sex = factor(sex),
         address = factor(address),
         famsize = factor(famsize),
         Pstatus = factor(Pstatus),
         schoolsup = factor(schoolsup),
         famsup = factor(famsup),
         paid = factor(paid),
         activities = factor(activities),
         nursery = factor(nursery),
         higher = factor(higher),
         internet = factor(internet),
         romantic = factor(romantic),
         reason = factor(reason))
```

student\_por

```
## # A tibble: 649 x 33
##   school sex   age address famsize Pstatus Medu Fedu Mjob  Fjob  reason
##   <fct> <fct> <dbl> <fct>   <fct>   <fct>   <dbl> <dbl> <chr>   <chr> <fct>
## 1 GP    F    18 U      GT3     A       4     4 at_home teach~ course
## 2 GP    F    17 U      GT3     T       1     1 at_home other  course
## 3 GP    F    15 U      LE3     T       1     1 at_home other  other
## 4 GP    F    15 U      GT3     T       4     2 health servi~ home
## 5 GP    F    16 U      GT3     T       3     3 other   other  home
## 6 GP    M    16 U      LE3     T       4     3 servic~ other  reputa~
## 7 GP    M    16 U      LE3     T       2     2 other   other  home
## 8 GP    F    17 U      GT3     A       4     4 other   teach~ home
## 9 GP    M    15 U      LE3     A       3     2 servic~ other  home
## 10 GP   M    15 U      GT3     T       3     4 other   other  home
## # ... with 639 more rows, and 22 more variables: guardian <chr>,
## #   traveltime <dbl>, studytime <dbl>, failures <dbl>, schoolsup <fct>,
## #   famsup <fct>, paid <fct>, activities <fct>, nursery <fct>, higher <fct>,
## #   internet <fct>, romantic <fct>, famrel <dbl>, freetime <dbl>, goout <dbl>,
## #   Dalc <dbl>, Walc <dbl>, health <dbl>, absences <dbl>, G1 <dbl>, G2 <dbl>,
## #   G3 <dbl>
```

### 3 EDA & Checking Assumptions

Before we begin our analysis, we wish to explore the distribution of the data and confirm it follows the typical assumptions of linear regression.


```
summary(student_por)
```

```
## school sex age address famsize Pstatus Medu
## GP:423 F:383 Min. :15.00 R:197 GT3:457 A: 80 Min. :0.000
## MS:226 M:266 1st Qu.:16.00 U:452 LE3:192 T:569 1st Qu.:2.000
## Median :17.00 Median :2.000
## Mean :16.74 Mean :2.515
## 3rd Qu.:18.00 3rd Qu.:4.000
## Max. :22.00 Max. :4.000
## Fedu Mjob Fjob reason
## Min. :0.000 Length:649 Length:649 course :285
## 1st Qu.:1.000 Class :character Class :character home :149
## Median :2.000 Mode :character Mode :character other : 72
## Mean :2.307 reputation:143
## 3rd Qu.:3.000
## Max. :4.000
## guardian traveltime studytime failures schoolsup
## Length:649 Min. :1.000 Min. :1.000 Min. :0.0000 no :581
## Class :character 1st Qu.:1.000 1st Qu.:1.000 1st Qu.:0.0000 yes: 68
## Mode :character Median :1.000 Median :2.000 Median :0.0000
## Mean :1.569 Mean :1.931 Mean :0.2219
## 3rd Qu.:2.000 3rd Qu.:2.000 3rd Qu.:0.0000
## Max. :4.000 Max. :4.000 Max. :3.0000
## famsup paid activities nursery higher internet romantic
## no :251 no :610 no :334 no :128 no : 69 no :151 no :410
## yes:398 yes: 39 yes:315 yes:521 yes:580 yes:498 yes:239
##
##
##
## famrel freetime goout Dalc Walc
## Min. :1.000 Min. :1.00 Min. :1.000 Min. :1.000 Min. :1.00
## 1st Qu.:4.000 1st Qu.:3.00 1st Qu.:2.000 1st Qu.:1.000 1st Qu.:1.00
## Median :4.000 Median :3.00 Median :3.000 Median :1.000 Median :2.00
## Mean :3.931 Mean :3.18 Mean :3.185 Mean :1.502 Mean :2.28
## 3rd Qu.:5.000 3rd Qu.:4.00 3rd Qu.:4.000 3rd Qu.:2.000 3rd Qu.:3.00
## Max. :5.000 Max. :5.00 Max. :5.000 Max. :5.000 Max. :5.00
## health absences G1 G2
## Min. :1.000 Min. : 0.000 Min. : 0.0 Min. : 0.00
## 1st Qu.:2.000 1st Qu.: 0.000 1st Qu.:10.0 1st Qu.:10.00
## Median :4.000 Median : 2.000 Median :11.0 Median :11.00
## Mean :3.536 Mean : 3.659 Mean :11.4 Mean :11.57
## 3rd Qu.:5.000 3rd Qu.: 6.000 3rd Qu.:13.0 3rd Qu.:13.00
## Max. :5.000 Max. :32.000 Max. :19.0 Max. :19.00
## G3
## Min. : 0.00
## 1st Qu.:10.00
## Median :12.00
```

```
## Mean   :11.91
## 3rd Qu.:14.00
## Max.   :19.00
```

```
student_por %>%
  ggplot(aes(x = G3)) +
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



The histogram above is skewed to the left when run our initial linear regression, we will check the residuals and QQ plots.

### 3.1 Running a simple linear regression

```
por_reg <- lm(G3 ~ ., data = student_por)
summary(por_reg)
```


```
##
## Call:
## lm(formula = G3 ~ ., data = student_por)
##
## Residuals:
```


```

##      Min      1Q  Median      3Q      Max
## -8.7618 -0.5148  0.0038  0.6047  5.4973
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.63823    0.96361   0.662 0.508011
## schoolMS      -0.19797    0.12783  -1.549 0.121992
## sexM          -0.12258    0.11778  -1.041 0.298423
## age           0.02869    0.04835   0.593 0.553208
## addressU       0.11446    0.12277   0.932 0.351565
## famsizeLE3     0.01560    0.11505   0.136 0.892197
## PstatusT      -0.09746    0.16256  -0.600 0.549055
## Medu          -0.09170    0.07097  -1.292 0.196799
## Fedu           0.04962    0.06461   0.768 0.442773
## Mjobhealth     0.26583    0.25225   1.054 0.292379
## Mjobother     -0.09351    0.14208  -0.658 0.510720
## Mjobservices   0.17255    0.17510   0.985 0.324808
## Mjobteacher    0.22115    0.23558   0.939 0.348232
## Fjobhealth    -0.44420    0.35256  -1.260 0.208189
## Fjobother     -0.33805    0.21391  -1.580 0.114544
## Fjobservices  -0.47121    0.22477  -2.096 0.036457 *
## Fjobteacher   -0.54368    0.31611  -1.720 0.085958 .
## reasonhome    -0.07885    0.13366  -0.590 0.555479
## reasonother   -0.36174    0.17236  -2.099 0.036251 *
## reasonreputation -0.16934    0.13990  -1.210 0.226584
## guardianmother -0.02513    0.12461  -0.202 0.840252
## guardianother  0.21732    0.24922   0.872 0.383539
## traveltime     0.13859    0.07459   1.858 0.063667 .
## studytime      0.04965    0.06620   0.750 0.453569
## failures      -0.25495    0.09900  -2.575 0.010254 *
## schoolsupyes   -0.18419    0.17319  -1.064 0.287969
## famsupyes      0.09456    0.10701   0.884 0.377230
## paidyes       -0.19166    0.21664  -0.885 0.376663
## activitiesyes  0.01208    0.10482   0.115 0.908275
## nurseryyes    -0.09562    0.12722  -0.752 0.452553
## higheryes      0.20749    0.18261   1.136 0.256285
## internetyes    0.08517    0.12955   0.657 0.511152
## romanticyes   -0.04209    0.10786  -0.390 0.696483
## famrel        -0.01597    0.05471  -0.292 0.770469
## freetime      -0.05002    0.05267  -0.950 0.342694
## goout         -0.01889    0.05041  -0.375 0.708033
## Dalc          -0.05194    0.07185  -0.723 0.469977
## Walc          -0.01693    0.05553  -0.305 0.760521
## health        -0.05522    0.03633  -1.520 0.129064
## absences      0.01359    0.01173   1.158 0.247198
## G1            0.12933    0.03762   3.438 0.000626 ***
## G2            0.87037    0.03495  24.906 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.249 on 607 degrees of freedom
## Multiple R-squared:  0.86, Adjusted R-squared:  0.8506
## F-statistic: 90.95 on 41 and 607 DF, p-value: < 2.2e-16


```


```
#par(mfrow = c(2, 2))
plot(por_reg)
```











From our results, we can conclude that a simple linear regression might not be the appropriate form of analysis for this data set. We will continue to explore cross validation techniques to compare its prediction accuracy to other models.

We should write out the model here

### 3.2 Validation Set Technique on Initial Regression


```
set.seed(1)
n <- nrow(student_por)
Z <- sample(n, .7*n)

reg.fit <- lm(G3 ~ ., data = student_por, subset = Z)
```

```
g3_predicted <- predict(reg.fit, student_por)
```

```
plot(student_por$G3[-Z], g3_predicted[-Z], xlab = "Actual Grades", ylab = "Predicted Grades", main = "P",
abline(0,1)
```

## Prediction Accuracy of Full Linear Model



```
mse_lm <- mean((student_por$G3 - g3_predicted)[-Z]^2)
```

A lot of the predictors do not appear to be significant, so we are going to use some variable selection methods to simplify the model while still maintaining accuracy. We hope to generate new models that have lower prediction MSE than this linear model.

### 3.3 Best Subset

```
# Takes a while to run
```

```
subsets <- regsubsets(G3 ~ ., data = student_por, nvmax = 15)
```

```
summary(subsets)
```

```
## Subset selection object
## Call: regsubsets.formula(G3 ~ ., data = student_por, nvmax = 15)
## 41 Variables (and intercept)
##               Forced in Forced out
## schoolMS      FALSE      FALSE
## sexM           FALSE      FALSE
## age            FALSE      FALSE
## addressU       FALSE      FALSE
## famsizeLE3     FALSE      FALSE
```

```

## PstatusT      FALSE      FALSE
## Medu          FALSE      FALSE
## Fedu          FALSE      FALSE
## Mjobhealth    FALSE      FALSE
## Mjobother     FALSE      FALSE
## Mjobservices  FALSE      FALSE
## Mjobteacher   FALSE      FALSE
## Fjobhealth    FALSE      FALSE
## Fjobother     FALSE      FALSE
## Fjobservices  FALSE      FALSE
## Fjobteacher   FALSE      FALSE
## reasonhome    FALSE      FALSE
## reasonother   FALSE      FALSE
## reasonreputation FALSE    FALSE
## guardianmother FALSE    FALSE
## guardianother FALSE    FALSE
## traveltime    FALSE      FALSE
## studytime     FALSE      FALSE
## failures      FALSE      FALSE
## schoolsupyes  FALSE      FALSE
## famsupyes     FALSE      FALSE
## paidyes       FALSE      FALSE
## activitiesyes FALSE      FALSE
## nurseryyes    FALSE      FALSE
## higheryes     FALSE      FALSE
## internetyes   FALSE      FALSE
## romanticyes   FALSE      FALSE
## famrel        FALSE      FALSE
## freetime      FALSE      FALSE
## goout         FALSE      FALSE
## Dalc          FALSE      FALSE
## Walc          FALSE      FALSE
## health        FALSE      FALSE
## absences      FALSE      FALSE
## G1            FALSE      FALSE
## G2            FALSE      FALSE
## 1 subsets of each size up to 15
## Selection Algorithm: exhaustive
##      schoolMS sexM age addressU famsizeLE3 PstatusT Medu Fedu Mjobhealth
## 1  ( 1 ) " "      " " " " " "      " "      " "      " "      " "
## 2  ( 1 ) " "      " " " " " "      " "      " "      " "      " "
## 3  ( 1 ) " "      " " " " " "      " "      " "      " "      " "
## 4  ( 1 ) " "      " " " " " "      " "      " "      " "      " "
## 5  ( 1 ) " "      "*" " " " "      " "      " "      " "      " "
## 6  ( 1 ) " "      "*" " " " "      " "      " "      " "      " "
## 7  ( 1 ) "*"      "*" " " " "      " "      " "      " "      " "
## 8  ( 1 ) "*"      "*" " " " "      " "      " "      " "      " "
## 9  ( 1 ) " "      "*" " " "*"      " "      " "      " "      " "
## 10 ( 1 ) " "      "*" " " "*"      " "      " "      " "      " "
## 11 ( 1 ) " "      "*" " " "*"      " "      " "      " "      " "
## 12 ( 1 ) " "      "*" " " "*"      " "      " "      " "      " "
## 13 ( 1 ) "*"      "*" " " "*"      " "      " "      " "      " "
## 14 ( 1 ) "*"      "*" " " "*"      " "      " "      " "      " "
## 15 ( 1 ) "*"      "*" " " "*"      " "      " "      " "      " "

```

##		Mjobother	Mjobservices	Mjobteacher	Fjobhealth	Fjobother	Fjobservices
## 1	( 1 )	" "	" "	" "	" "	" "	" "
## 2	( 1 )	" "	" "	" "	" "	" "	" "
## 3	( 1 )	" "	" "	" "	" "	" "	" "
## 4	( 1 )	" "	" "	" "	" "	" "	" "
## 5	( 1 )	" "	" "	" "	" "	" "	" "
## 6	( 1 )	" "	" "	" "	" "	" "	" "
## 7	( 1 )	" "	" "	" "	" "	" "	" "
## 8	( 1 )	"*"	" "	" "	" "	" "	" "
## 9	( 1 )	"*"	" "	" "	" "	" "	" "
## 10	( 1 )	"*"	" "	" "	" "	" "	" "
## 11	( 1 )	"*"	" "	" "	" "	" "	" "
## 12	( 1 )	"*"	" "	" "	" "	" "	" "
## 13	( 1 )	"*"	" "	" "	" "	" "	" "
## 14	( 1 )	"*"	" "	" "	" "	" "	" "
## 15	( 1 )	"*"	" "	" "	" "	" "	" "

##		Fjobteacher	reasonhome	reasonother	reasonreputation	guardianmother
## 1	( 1 )	" "	" "	" "	" "	" "
## 2	( 1 )	" "	" "	" "	" "	" "
## 3	( 1 )	" "	" "	"*"	" "	" "
## 4	( 1 )	" "	" "	"*"	" "	" "
## 5	( 1 )	" "	" "	"*"	" "	" "
## 6	( 1 )	" "	" "	"*"	" "	" "
## 7	( 1 )	" "	" "	"*"	" "	" "
## 8	( 1 )	" "	" "	"*"	" "	" "
## 9	( 1 )	" "	" "	"*"	" "	" "
## 10	( 1 )	" "	" "	"*"	" "	" "
## 11	( 1 )	" "	" "	"*"	" "	" "
## 12	( 1 )	" "	" "	"*"	" "	" "
## 13	( 1 )	" "	" "	"*"	" "	" "
## 14	( 1 )	" "	" "	"*"	" "	" "
## 15	( 1 )	" "	" "	"*"	" "	" "

##		guardianother	traveltime	studytime	failures	schoolsupyes	famsupyes
## 1	( 1 )	" "	" "	" "	" "	" "	" "
## 2	( 1 )	" "	" "	" "	" "	" "	" "
## 3	( 1 )	" "	" "	" "	" "	" "	" "
## 4	( 1 )	" "	" "	" "	"*"	" "	" "
## 5	( 1 )	" "	" "	" "	"*"	" "	" "
## 6	( 1 )	" "	" "	" "	"*"	" "	" "
## 7	( 1 )	" "	"*"	" "	"*"	" "	" "
## 8	( 1 )	" "	"*"	" "	"*"	" "	" "
## 9	( 1 )	" "	"*"	" "	"*"	" "	" "
## 10	( 1 )	"*"	"*"	" "	"*"	" "	" "
## 11	( 1 )	"*"	"*"	" "	"*"	" "	" "
## 12	( 1 )	"*"	"*"	" "	"*"	" "	" "
## 13	( 1 )	"*"	"*"	" "	"*"	" "	" "
## 14	( 1 )	"*"	"*"	" "	"*"	" "	" "
## 15	( 1 )	"*"	"*"	" "	"*"	"*"	" "

##		paidyes	activitiesyes	nurseryyes	higheryes	internetyes	romanticyes
## 1	( 1 )	" "	" "	" "	" "	" "	" "
## 2	( 1 )	" "	" "	" "	" "	" "	" "
## 3	( 1 )	" "	" "	" "	" "	" "	" "
## 4	( 1 )	" "	" "	" "	" "	" "	" "
## 5	( 1 )	" "	" "	" "	" "	" "	" "

```
## 6 ( 1 ) " " " " " " " " " "
## 7 ( 1 ) " " " " " " " " " "
## 8 ( 1 ) " " " " " " " " " "
## 9 ( 1 ) " " " " " " " " " "
## 10 ( 1 ) " " " " " " " " " "
## 11 ( 1 ) " " " " " " " " " "
## 12 ( 1 ) " " " " " " " " " "
## 13 ( 1 ) " " " " " " " " " "
## 14 ( 1 ) " " " " " " "*" " " "
## 15 ( 1 ) " " " " " " "*" " " "

##      famrel freetime goout Dalc Walc health absences G1 G2
## 1 ( 1 ) " " " " " " " " " " " " "*"
## 2 ( 1 ) " " " " " " " " " " " " "*" "*"
## 3 ( 1 ) " " " " " " " " " " " " "*" "*"
## 4 ( 1 ) " " " " " " " " " " " " "*" "*"
## 5 ( 1 ) " " " " " " " " " " " " "*" "*"
## 6 ( 1 ) " " " " " " " " " " "*" "*" "*"
## 7 ( 1 ) " " " " " " " " " " " " "*" "*"
## 8 ( 1 ) " " " " " " " " " " " " "*" "*"
## 9 ( 1 ) " " " " " " " " " " "*" "*" "*"
## 10 ( 1 ) " " " " " " " " " " "*" "*" "*"
## 11 ( 1 ) " " " " " " "*" " " " "*" "*" "*"
## 12 ( 1 ) " " " " " " "*" " " "*" "*" "*"
## 13 ( 1 ) " " " " " " "*" " " "*" "*" "*"
## 14 ( 1 ) " " " " " " "*" " " "*" "*" "*"
## 15 ( 1 ) " " " " " " "*" " " "*" "*" "*"

```

```
summary(subsets)$adjr2
```

```
## [1] 0.8434889 0.8472902 0.8489562 0.8501271 0.8507762 0.8513288 0.8518147
## [8] 0.8522527 0.8526228 0.8528667 0.8531441 0.8533129 0.8534618 0.8534808
## [15] 0.8535407

```

```
summary(subsets)$cp
```

```
## [1] 32.5841845 17.1053566 10.8932613 6.8369620 5.0413803 3.6686368
## [7] 2.5898385 1.7227267 1.1515489 1.1240570 0.9573861 1.2563559
## [13] 1.6420508 2.5809380 3.3465888

```

```
summary(subsets)$bic
```

```
## [1] -1191.705 -1202.191 -1203.840 -1203.422 -1200.772 -1197.715 -1194.376
## [8] -1190.835 -1187.002 -1182.618 -1178.385 -1173.676 -1168.881 -1163.512
## [15] -1158.327

```

### 3.4 Set validation for Best Subset

### 3.5 Best Subset

```
which.max(summary(subsets)$adjr2)
```

```
## [1] 15
```

```
which.min(abs(summary(subsets)$cp - 1:15))
```

```
## [1] 5
```

```
which.min(summary(subsets)$bic)
```

```
## [1] 3
```


Didn't use adjr2 model bc they all make little difference.

### 3.6 Model Based on Mallows's Cp

```
reg.bestsubCP <- lm(G3 ~ sex + reason + failures + G1 + G2, data = student_por, subset = Z)
```

```
g3_pred_bestsubCP <- predict(reg.bestsubCP, student_por)
```

```
plot(student_por$G3[-Z], g3_pred_bestsubCP[-Z], xlab = "Actual Grades", ylab = "Predicted Grades", main = "Predicted vs. Actual Grades of Reduced Model Based on Cp", abline(0,1))
```




```
mse_cp <- mean((student_por$G3 - g3_pred_bestsubCP)[-Z] ^ 2)
```

### 3.7 Model Based on BIC

```
reg.bestsubBIC <- lm(G3 ~ reason + G1 + G2, data = student_por, subset = Z)
```

```
g3_pred_bestsubBIC <- predict(reg.bestsubBIC, student_por)
```

```
plot(student_por$G3[-Z], g3_pred_bestsubBIC[-Z], xlab = "Actual Grades", ylab = "Predicted Grades", main = "Predicted vs. Actual Grades of Reduced Model Based on BIC", abline(0,1))
```



```
mse_bic <- mean((student_por$G3 - g3_pred_bestsubBIC)[-Z] ^ 2)
```

### 3.8 Step Functions

```
summary(forward)
```

```
##
## Call:
## lm(formula = G3 ~ G2 + G1 + failures + reason + absences + sex +
```



```
##      school + traveltime + health, data = student_por)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -9.0833 -0.5178 -0.0053  0.6398  5.2097
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.44063    0.34169   1.290 0.197678
## G2              0.87996    0.03379  26.042 < 2e-16 ***
## G1              0.13706    0.03615   3.792 0.000164 ***
## failures       -0.24049    0.09074  -2.650 0.008244 **
## reasonhome     -0.09222    0.13010  -0.709 0.478659
## reasonother    -0.44994    0.16627  -2.706 0.006990 **
## reasonreputation -0.16537    0.13290  -1.244 0.213816
## absences        0.01623    0.01100   1.476 0.140522
## sexM           -0.20022    0.10191  -1.965 0.049894 *
## schoolMS       -0.22981    0.11621  -1.977 0.048419 *
## traveltime      0.11228    0.06839   1.642 0.101138
## health         -0.05394    0.03469  -1.555 0.120451
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.242 on 637 degrees of freedom
## Multiple R-squared:  0.8547, Adjusted R-squared:  0.8522
## F-statistic: 340.7 on 11 and 637 DF,  p-value: < 2.2e-16
```

```
summary(backward)
```

```
##
## Call:
## lm(formula = G3 ~ school + sex + reason + traveltime + failures +
##      health + absences + G1 + G2, data = student_por)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -9.0833 -0.5178 -0.0053  0.6398  5.2097
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.44063    0.34169   1.290 0.197678
## schoolMS       -0.22981    0.11621  -1.977 0.048419 *
## sexM           -0.20022    0.10191  -1.965 0.049894 *
## reasonhome     -0.09222    0.13010  -0.709 0.478659
## reasonother    -0.44994    0.16627  -2.706 0.006990 **
## reasonreputation -0.16537    0.13290  -1.244 0.213816
## traveltime      0.11228    0.06839   1.642 0.101138
## failures       -0.24049    0.09074  -2.650 0.008244 **
## health         -0.05394    0.03469  -1.555 0.120451
## absences        0.01623    0.01100   1.476 0.140522
## G1              0.13706    0.03615   3.792 0.000164 ***
## G2              0.87996    0.03379  26.042 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
##
## Residual standard error: 1.242 on 637 degrees of freedom
## Multiple R-squared:  0.8547, Adjusted R-squared:  0.8522
## F-statistic: 340.7 on 11 and 637 DF,  p-value: < 2.2e-16
```

Forward and backward step functions yield the exact same model; proceeding with forward step-generated model.

### 3.9 Set validation

```
reg.forward <- lm(G3 ~ G2 + G1 + failures + reason + absences + sex + school + traveltime + health, data=student_por)
g3_pred_forward <- predict(reg.forward, student_por)
```

```
plot(student_por$G3[-Z], g3_pred_forward[-Z])
abline(0, 1)
```



```
mse_valSet <- mean((student_por$G3 - g3_pred_forward)[-Z] ^ 2)
```

### 3.10 Ridge Regression & LASSO Preparation

```

# Training/test split
# set.seed(1)
# train <- sample(1:n, n/2)
G3_test <- student_por$G3[-Z]

# Creating model matrix for rr and lasso calculations
x_col <- model.matrix(G3 ~ ., student_por)[, -1]

```

### 3.11 Ridge Regression

```

set.seed(1)
cv.out1 <- cv.glmnet(x_col, student_por$G3, alpha = 0) # alpha = 0 ---> Ridge regression
predict(cv.out1, s = cv.out1$lambda.min, type = "coefficients")

```

```

## 42 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  0.647281475
## schoolMS    -0.203906036
## sexM        -0.140157325
## age         0.052338494
## addressU     0.134471281
## famsizeLE3   0.031162924
## PstatusT    -0.065759455
## Medu        -0.052585350
## Fedu         0.042315840
## Mjobhealth   0.214144623
## Mjobother    -0.125812519
## Mjobservices 0.107348800
## Mjobteacher  0.126800272
## Fjobhealth   -0.238572263
## Fjobother    -0.156282858
## Fjobservices -0.300920468
## Fjobteacher  -0.249881588
## reasonhome   -0.066205753
## reasonother  -0.362017744
## reasonreputation -0.112276025
## guardianmother -0.032801219
## guardianother 0.212481561
## traveltime   0.111686852
## studytime    0.060781281
## failures     -0.316944745
## schoolsupyes -0.201688220
## famsupyes    0.087520158
## paidyes      -0.153298259
## activitiesyes 0.018965743
## nurseryyes   -0.084879671
## higheryes    0.280580244
## internetyes  0.099020055
## romanticyes  -0.087129209
## famrel       0.008539498

```

```
## freetime      -0.053552694
## goout         -0.025522281
## Dalc          -0.056693604
## Walc          -0.025883700
## health        -0.065912944
## absences      0.011871609
## G1            0.258399918
## G2            0.683417827
```

```
rr.mod <- glmnet(x_col[Z, ], student_por$G3[Z], alpha = 0, lambda = cv.out1$lambda.min)
rr.pred <- predict(rr.mod, s = cv.out1$lambda.min, newx = x_col[-Z, ])

mse_rr <- mean((rr.pred - student_por$G3[-Z])^2)
```

$\lambda = .30$

### 3.12 LASSO

```
set.seed(1)
cv.out2 <- cv.glmnet(x_col, student_por$G3, alpha = 1)
predict(cv.out2, s = cv.out2$lambda.min, type = "coefficients")
```

```
## 42 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  0.46985582
## schoolMS    -0.03190401
## sexM        -0.01841156
## age         .
## addressU    .
## famsizeLE3  .
## PstatusT    .
## Medu        .
## Fedu        .
## Mjobhealth   .
## Mjobother    .
## Mjobservices .
## Mjobteacher  .
## Fjobhealth   .
## Fjobother    .
## Fjobservices .
## Fjobteacher  .
## reasonhome   .
## reasonother  -0.14557639
## reasonreputation .
## guardianmother .
## guardianother .
## traveltime   .
## studytime    .
## failures     -0.09120067
## schoolsupyes .
## famsupyes    .
```

```
## paidyes      .
## activitiesyes .
## nurseryyes   .
## higheryes    .
## internetyes  .
## romanticyes  .
## famrel       .
## freetime     .
## goout        .
## Dalc         .
## Walc         .
## health       .
## absences     .
## G1           0.12252007
## G2           0.87247067
```

```
 $\lambda = .10$ 
```

```
lasso.mod <- glmnet(x_col[Z, ], student_por$G3[Z], alpha = 1, lambda = cv.out2$lambda.min)
lasso.pred <- predict(lasso.mod, s = cv.out2$lambda.min, newx = x_col[-Z, ])

mse_lasso <- mean((lasso.pred - student_por$G3[-Z])^2)
```

```
student_por.dimred <- lm(G3 ~ school + sex + reason + failures + G1 + G2, student_por)
summary(student_por.dimred)
```

```
##
## Call:
## lm(formula = G3 ~ school + sex + reason + failures + G1 + G2,
##     data = student_por)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.2349 -0.4970  0.0057  0.6422  5.3485
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.53906    0.27382   1.969 0.049420 *
## schoolMS      -0.20979    0.11165  -1.879 0.060691 .
## sexM          -0.21513    0.10121  -2.126 0.033927 *
## reasonhome    -0.08152    0.12873  -0.633 0.526772
## reasonother   -0.45325    0.16677  -2.718 0.006748 **
## reasonreputation -0.14147    0.13181  -1.073 0.283529
## failures      -0.22420    0.09075  -2.471 0.013751 *
## G1             0.12912    0.03608   3.579 0.000371 ***
## G2             0.88212    0.03382  26.081 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.247 on 640 degrees of freedom
## Multiple R-squared:  0.8529, Adjusted R-squared:  0.8511
## F-statistic: 463.9 on 8 and 640 DF, p-value: < 2.2e-16
```

### 3.13 Comparing MSEs

```
tibble("method" = c("BIC-Minimized", "Cp-Minimized", "LASSO", "Linear Regression", "Ridge Regression",  
  "MSE" = c(mse_bic, mse_cp, mse_lasso, mse_lm, mse_rr, mse_valSet)) %>%  
  arrange(MSE)
```

```
## # A tibble: 6 x 2  
##   method      MSE  
##   <chr>      <dbl>  
## 1 BIC-Minimized    1.43  
## 2 AIC-Minimized    1.46  
## 3 Cp-Minimized     1.47  
## 4 LASSO            1.53  
## 5 Linear Regression 1.55  
## 6 Ridge Regression  1.60
```

```
reg.bestsubBIC
```

```
##  
## Call:  
## lm(formula = G3 ~ reason + G1 + G2, data = student_por, subset = Z)  
##  
## Coefficients:  
##      (Intercept)      reasonhome      reasonother      reasonreputation  
##      -0.01247      -0.08366      -0.33815      -0.10231  
##           G1           G2  
##      0.10995      0.92593
```

```
reg.forward # Picking this one
```

```
##  
## Call:  
## lm(formula = G3 ~ G2 + G1 + failures + reason + absences + sex +  
##      school + traveltime + health, data = student_por, subset = Z)  
##  
## Coefficients:  
##      (Intercept)           G2           G1      failures  
##      0.50898      0.90144      0.09950     -0.37556  
##      reasonhome      reasonother      reasonreputation      absences  
##      -0.13167      -0.33802      -0.21118      0.01524  
##           sexM      schoolMS      traveltime      health  
##      -0.20738      -0.22059      0.16919     -0.04444
```

```
reg.bestsubCP
```

```
##  
## Call:  
## lm(formula = G3 ~ sex + reason + failures + G1 + G2, data = student_por,  
##      subset = Z)  
##
```

```
## Coefficients:
##      (Intercept)          sexM          reasonhome          reasonother
##      0.55807         -0.18604         -0.11971         -0.36133
## reasonreputation      failures              G1              G2
##      -0.14844         -0.35971          0.09759          0.90438
```

Picking forward-selected candidate model b/c best balance of number of predictors while sacrificing only a little accuracy.

```
summary(reg.forward)
```

```
##
## Call:
## lm(formula = G3 ~ G2 + G1 + failures + reason + absences + sex +
##      school + traveltime + health, data = student_por, subset = Z)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.8824 -0.4672 -0.0923  0.6427  5.0271
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.50898    0.42147   1.208 0.227839
## G2              0.90144    0.03983  22.634 < 2e-16 ***
## G1              0.09950    0.04334   2.296 0.022162 *
## failures       -0.37556    0.10959  -3.427 0.000667 ***
## reasonhome     -0.13167    0.15640  -0.842 0.400327
## reasonother    -0.33802    0.20922  -1.616 0.106879
## reasonreputation -0.21118    0.16202  -1.303 0.193108
## absences        0.01524    0.01385   1.100 0.271973
## sexM           -0.20738    0.12313  -1.684 0.092851 .
## schoolMS       -0.22059    0.13985  -1.577 0.115432
## traveltime      0.16919    0.08337   2.029 0.043026 *
## health         -0.04444    0.04150  -1.071 0.284835
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.264 on 442 degrees of freedom
## Multiple R-squared:  0.854, Adjusted R-squared:  0.8504
## F-statistic: 235.1 on 11 and 442 DF, p-value: < 2.2e-16
```

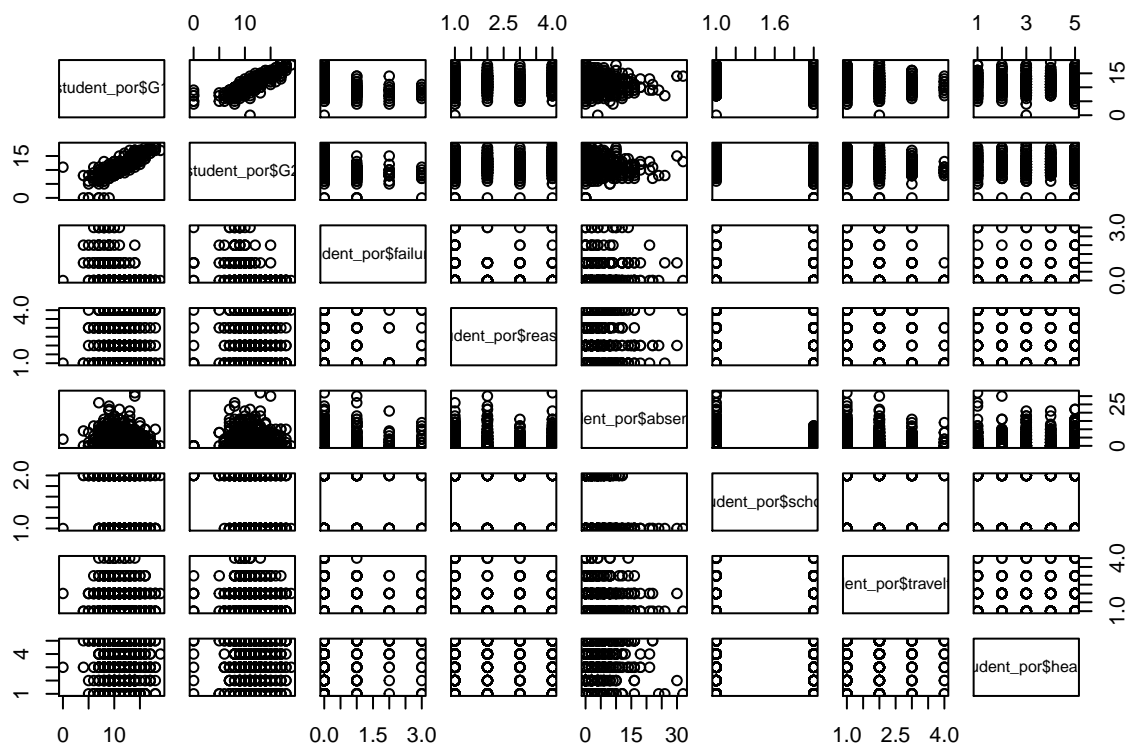
```
summary(
  lm(formula = G3 ~ G2 + G1 + failures + reason + absences + sex +
      school + traveltime + health, data = student_por)
)
```

```
##
## Call:
## lm(formula = G3 ~ G2 + G1 + failures + reason + absences + sex +
##      school + traveltime + health, data = student_por)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -9.0833 -0.5178 -0.0053 0.6398 5.2097
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.44063    0.34169   1.290 0.197678
## G2              0.87996    0.03379  26.042 < 2e-16 ***
## G1              0.13706    0.03615   3.792 0.000164 ***
## failures       -0.24049    0.09074  -2.650 0.008244 **
## reasonhome     -0.09222    0.13010  -0.709 0.478659
## reasonother    -0.44994    0.16627  -2.706 0.006990 **
## reasonreputation -0.16537    0.13290  -1.244 0.213816
## absences        0.01623    0.01100   1.476 0.140522
## sexM           -0.20022    0.10191  -1.965 0.049894 *
## schoolMS       -0.22981    0.11621  -1.977 0.048419 *
## traveltime      0.11228    0.06839   1.642 0.101138
## health         -0.05394    0.03469  -1.555 0.120451
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.242 on 637 degrees of freedom
## Multiple R-squared:  0.8547, Adjusted R-squared:  0.8522
## F-statistic: 340.7 on 11 and 637 DF, p-value: < 2.2e-16
```

```
pairs(tibble(student_por$G1,
              student_por$G2,
              student_por$failures,
              student_por$reason,
              student_por$absences,
              student_por$school,
              student_por$traveltime,
              student_por$health))
```





```
cor(data.frame(student_por$G1, student_por$G2))
```

```
##           student_por.G1 student_por.G2
## student_por.G1      1.000000      0.8649816
## student_por.G2      0.8649816      1.0000000
```

```
car::vif(reg.forward)
```

```
##           GVIF Df GVIF^(1/(2*Df))
## G2           3.947302 1          1.986782
## G1           3.931017 1          1.982679
## failures     1.233508 1          1.110634
## reason       1.175917 3          1.027376
## absences     1.082722 1          1.040539
## sex          1.044820 1          1.022164
## school       1.264176 1          1.124356
## traveltime   1.101912 1          1.049720
## health       1.066354 1          1.032644
```

```
reg.forward_mod <- lm(G3 ~ G2 + failures + reason + absences + sex + school + traveltime + health, stud
```

```
summary(reg.forward_mod)
```

```
##
```

```
## Call:
## lm(formula = G3 ~ G2 + failures + reason + absences + sex + school +
##      traveltime + health, data = student_por)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.0375 -0.4999 -0.0428  0.6332  5.1354
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.81202    0.33077   2.455  0.01436 *
## G2              0.98521    0.01947  50.606 < 2e-16 ***
## failures       -0.27267    0.09129  -2.987  0.00293 **
## reasonhome     -0.07680    0.13139  -0.584  0.55910
## reasonother    -0.44246    0.16799  -2.634  0.00865 **
## reasonreputation -0.15001    0.13422  -1.118  0.26415
## absences        0.01204    0.01105   1.089  0.27641
## sexM           -0.21827    0.10286  -2.122  0.03422 *
## schoolMS       -0.28698    0.11643  -2.465  0.01397 *
## traveltime      0.11205    0.06911   1.621  0.10543
## health         -0.04955    0.03503  -1.414  0.15775
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.255 on 638 degrees of freedom
## Multiple R-squared:  0.8514, Adjusted R-squared:  0.8491
## F-statistic: 365.7 on 10 and 638 DF, p-value: < 2.2e-16
```

```
car::vif(reg.forward_mod)
```

```
##              GVIF Df GVIF^(1/(2*Df))
## G2           1.324014 1      1.150658
## failures     1.206829 1      1.098558
## reason       1.190013 3      1.029418
## absences     1.083037 1      1.040691
## sex          1.054733 1      1.027002
## school       1.268032 1      1.126069
## traveltime   1.101477 1      1.049513
## health       1.056308 1      1.027768
```